

# System Programming

## *3. File IO (2): System Call*

Yunseok Rhee

[rheey@hufs.ac.kr](mailto:rheey@hufs.ac.kr)

Division of Computer Engineering



# Linux File System (1)

- Each file in a file system has its own ***inode***
- An inode is a data structure having all information on a file.
- inodes of all files reside in a ***disk***
- inode contents (C struct)
  - file name
  - file type (regular, directory,...)
  - file owner id
  - access permission  
rwxr-xr-x (for owner, group, others)
  - creation/modified time
  - file size
  - file data block addr. table (see the next page!)
  - ...

# Linux File System (2)

## ■ File types

- regular file
- directory file
- FIFO file (pipe)
- special files (IO devices)
- symbolic link files

A “john” directory file

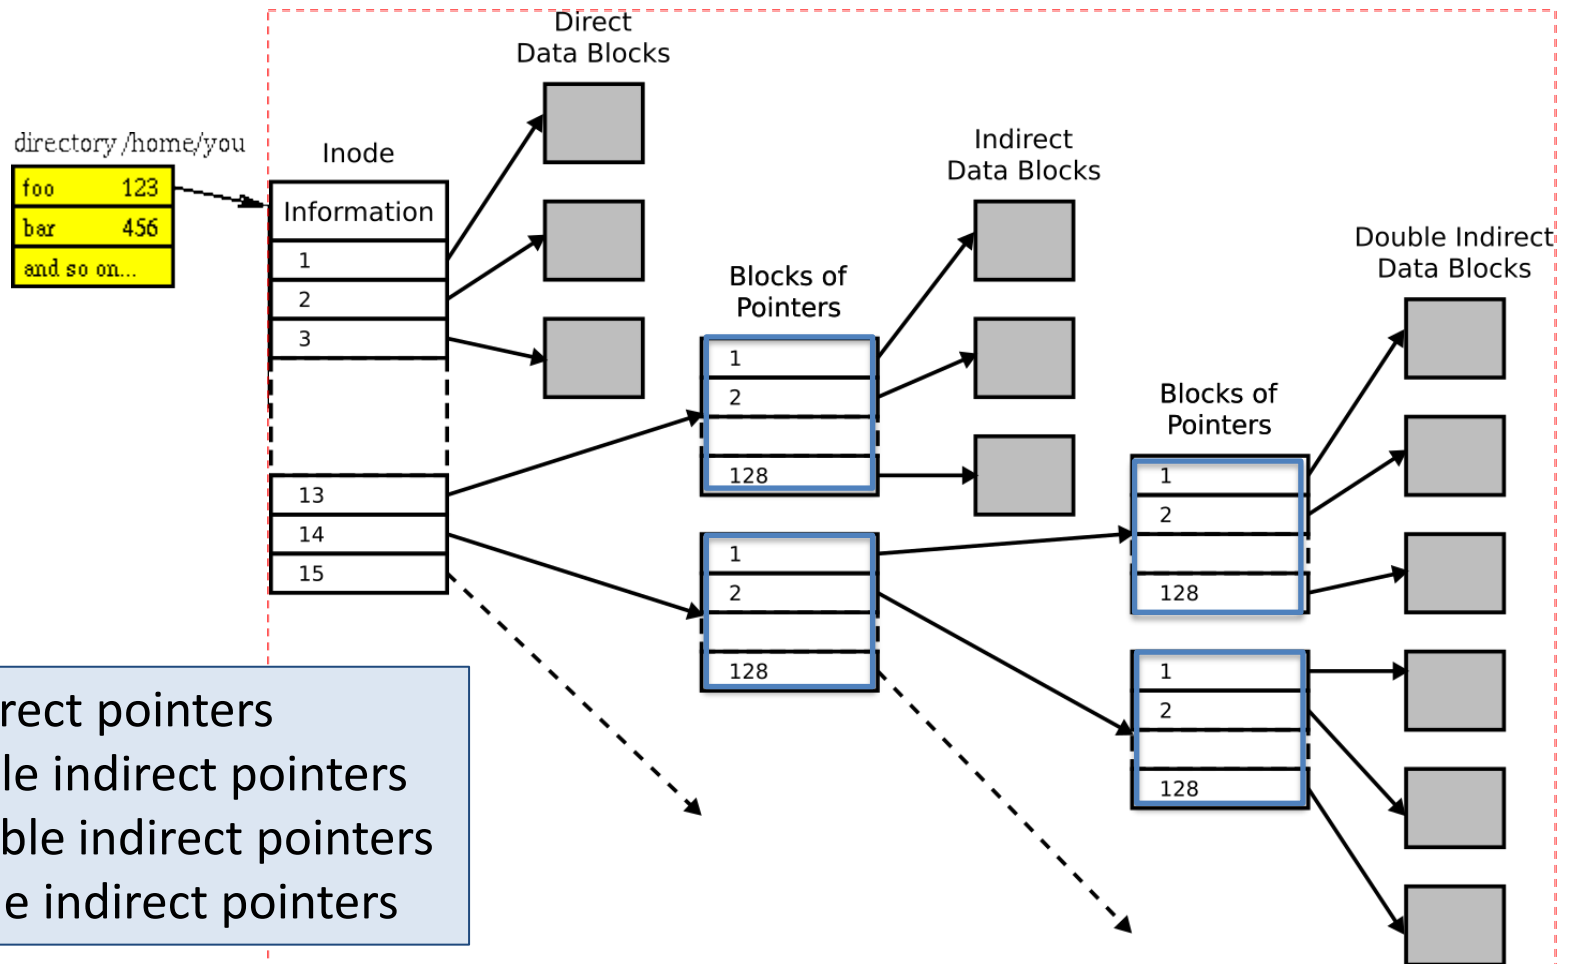
i	. (john)
j	.. (parent)
k	File name A
l	File name B
m	File name C

## ■ Directory file

- A directory is just a file whose content is the list of (inode #, file name) in the same directory.
- inode is a data structure which contains all the information about the file and file data blocks
- inode # is a unique file id number in the file system
- “ls -al john” is a shell command that just displays the “john” directory file

# Inode structure example

- block size: 512 byte, block pointer: 4 bytes

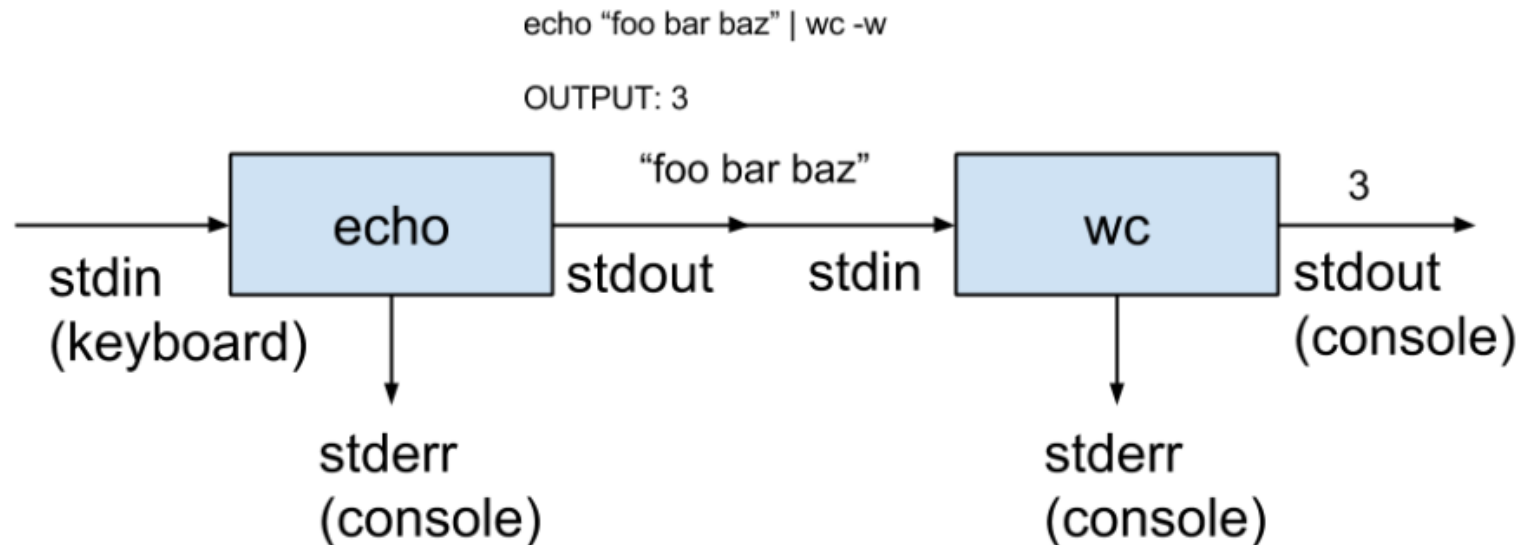


# Linux File Types

- Ordinary File (Regular File)
  - Text, binary files
- Directory File
  - A file that includes the set of (*file-name*, *inode #*) of the directory.
- Character Special File
  - Character-oriented device (e.g. Keyboard)
- Block Special File
  - Block-oriented device (e.g. HDD file systems, eth0 )
- FIFO file
  - Named *pipe* / *Unnamed pipe*  
*cf. pipe in a process is usually unnamed.*
- Symbolic link file
  - a file which points to another file  
*cf. **hardlink** is NOT a file.*

# File Descriptor (1)

- A *file descriptor* (or *file handle*) is a small, non-negative integer which identifies a file to the kernel.
  - Traditionally, `stdin`, `stdout` and `stderr` are 0, 1 and 2 respectively.



- Relying on “magic numbers” is BAD.
  - Use `STDIN_FILENO`, `STDOUT_FILENO` and `STDERR_FILENO` defined in `unistd.h` or `stdin`, `stdout`, and `stderr` defined in `stdio.h`.

# File Descriptor (2)

- Maximum number of files
  - a process can open 1024 files
  - we can check the system resource configuration

```
$ ulimit -a
core file size      (blocks, -c) 0
data seg size       (kbytes, -d) unlimited
scheduling priority (-e) 0
file size           (blocks, -f) unlimited
pending signals      (-i) 194273
max locked memory    (kbytes, -l) 64
max memory size      (kbytes, -m) unlimited
open files           (-n) 1024
pipe size            (512 bytes, -p) 8
.....
```

# Basic File I/Os

- 5 fundamental Unix/Linux file I/Os
  - `open(2)`
  - `close(2)`
  - `lseek(2)`
  - `read(2)`
  - `write(2)`



# File open (1)

```
#include <fcntl.h>
```

```
int open(const char *path, int oflag);
```

```
int open(const char *path, int oflag, mode_t mode);
```

## ■ parameters

- *path*: name of the file to open or create
- *oflag*: file open options
- *mode*: access permission (at file creation)

## ■ return

- *file descriptor* if OK
- -1 on error

# File open (2)

- *oflag* options

- must be one of these

option1	meaning	<fcntl.h> defined
O_RDONLY	open for reading only	0
O_WRONLY	open for writing only	1
O_RDWR	open for reading & writing	2

- and can be OR'ed with any of these (by "|")

option2	meaning
O_CREAT	create a file if the file does not exist.
O_EXCL	used with O_CREAT, return an error if the file already exists.
O_TRUNC	if the file exists, make it empty.
O_APPEND	write from the end of the file.
O_SYNC	do disk synchronization when does file I/O.

# File access modes (1)

mode	meaning
S_ISUID	set-user-id at execution
S_ISGID	set-group-id at execution
S_ISVTX	set sticky bit
S_IRWXU	owner RWX
S_IRUSR	owner R
S_IWUSR	owner W
S_IXUSR	owner X
S_IRWXG	group RWX
S_IRGRP	group R
S_IWGRP	group W
S_IXGRP	group X
S_IRWXO	others RWX
S_IROTH	others R
S_IWOTH	others W
S_IXOTH	others X



will be explained later

# File close

```
#include <unistd.h>

int close(int fd);
```

- parameters
  - *fd*: file descriptor
- return
  - 0 if OK
  - -1 on error

# File open example

## *open-ex.c*

```
#include <fcntl.h>

int main(int argc, char * argv[])
{
    FILE *fpo; // file pointer
    int fdo;    // file descriptor

    if(argc != 2) {
        perror(argv[0]);
        return 1;
    }
    if((fdo = open(argv[1], O_RDWR | O_CREAT | O_TRUNC,
                    S_IRUSR | S_IWUSR)) == -1) {
        perror(argv[1]);
        return 1;
    }
    if((fpo = fdopen(fdo, "r+")) == NULL) {
        perror("fdopen");
        return 2;
    }
    fprintf(fpo, "Hello, world! \n");
    fclose(fpo);
}
```

```
./a.out test.txt
$ cat test.txt
Hello, world!
$
```

# File creation

```
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

int creat(const char *path, mode_t mode);
```

## ■ parameters

- *path* : file path name
- *mode* : access permission

## ■ return

- *file descriptor* if OK
- -1 on error

# File seeking

```
#include <sys/types.h>
#include <unistd.h>

off_t lseek(int fd, off_t offset, int whence);
```

## ■ parameters

- *fd* : file descriptor
- *offset* : offset from the beginning to seek (move)
- *whence* : SEEK\_SET, SEEK\_CUR, SEEK\_END

## ■ return

- new offset value if OK
- -1 on error

# File create/lseek example

## *create-ex.c*

```
#include <sys/types.h>
#include <sys/stat.h>
#include <sys/fcntl.h>
#include <unistd.h>

int main(void)
{
    int fd;
    char buf1[] = "Test1 data";
    char buf2[] = "Test2 data";

    if ((fd == creat ("test.txt", S_IRUSR | S_IWUSR | S_IRGRP |
                      S_IROTH)) < 0) {
        printf("creat error");
        return 1;
    }
    write(fd, buf1, 10);
    if(lseek(fd, 6L, SEEK_SET) == -1) {
        printf("lseek error");
        return 2;
    }
    write(fd, buf2, 10);

    return 0;
}
```

```
$ ls
a.out  test.c
$ ./a.out
$ ls
a.out  test.c  test.txt
$ cat  test.txt
Test1 Test2 data
$
```



# File reading

```
#include <unistd.h>
```

```
ssize_t read(int fd, void *buf, size_t nbyte);
```

## ■ parameters

- *fd* : file descriptor
- *buf* : buffer address
- *nbyte* : number of bytes to read

## ■ return

- number of bytes read successfully if OK
- -1 on error

# File writing

```
#include <unistd.h>
```

```
ssize_t write(int fd, const void *buf, size_t nbyte);
```

## ■ parameters

- *fd* : file descriptor
- *buf* : buffer address
- *nbyte* : number of bytes to write

## ■ return

- number of bytes written successfully if OK
- -1 on error

# File copy example (1)

## *fcopy2-ex.c*

```
#include <stdio.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

#define BUFFER_SIZE 1024

int main(int argc, char *argv[])
{
    int fdi, fdo;
    char buf[BUFFER_SIZE];
    ssize_t n;

    if(argc != 3) {
        perror(argv[0]);
        return 1;
    }
```

# File copy example (2)

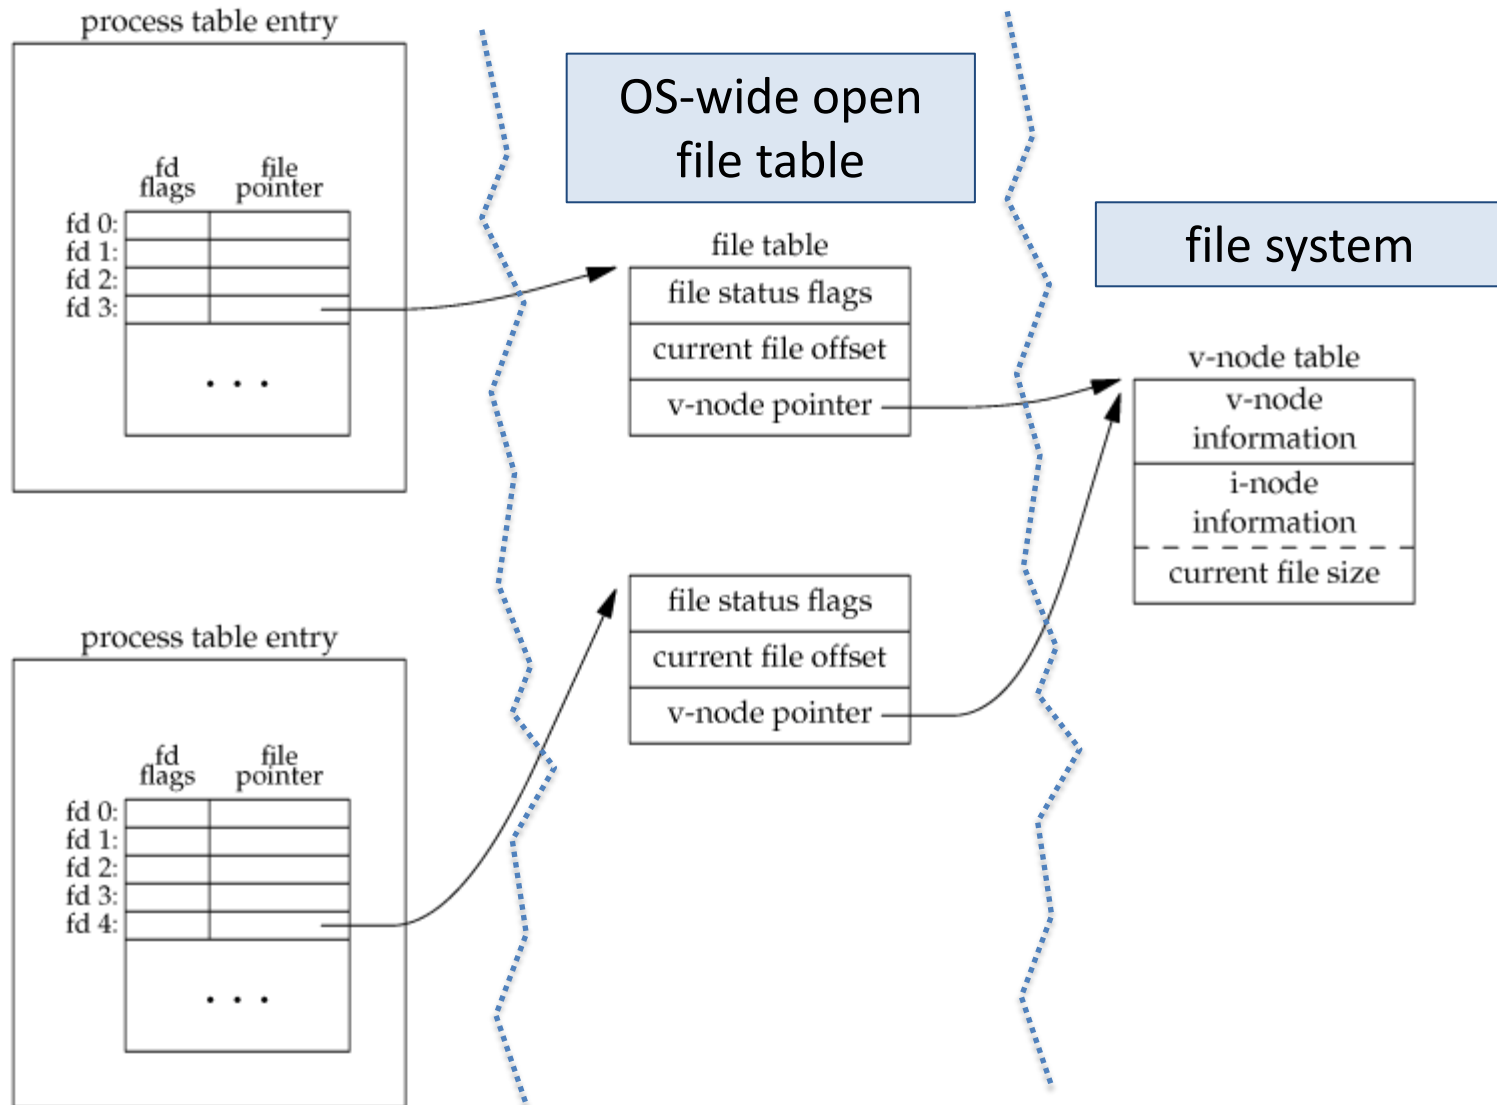
*fcopy2-ex.c*

```
    if((fdi = open(argv[1], O_RDONLY)) == -1) {
        perror(argv[1]);
        return 2;
    }

    if((fdo = open(argv[2], O_WRONLY | O_CREAT | O_EXCL,
                    S_IRUSR | S_IWUSR)) == -1) {
        perror(argv[2]);
        return 3;
    }
    while((n = read(fdi, buf, BUFFER_SIZE)) > 0)
        write(fdo, buf, n);

    close(fdi);
    close(fdo);
    return 0;
}
```


# File descriptors & File table



# Duplication of file descriptor (1)

```
#include <unistd.h>
int dup(int fd);
```

- parameters
  - *fd* : file descriptor to duplicate
- return
  - newly duplicated file descriptor if OK
  - -1 on error



# Duplication of file descriptor (2)

```
#include <unistd.h>
int dup2(int fd1, int fd2);
```

- parameters
  - *fd1* : source file descriptor
  - *fd2* : destination file descriptor
- return
  - copied file descriptor (should be same as *fd2*) if OK
  - -1 on error
- note
  - functionally same as *dup* except that *dup2* designates destination file descriptor (*fd2*) the user wants

# I/O redirection example (1)

*io-redir.c*

```
#include <stdio.h>
#include <unistd.h>
#include <sys/types.h>
#include <sys/stat.h>
#include <fcntl.h>

int main()
{
    int backup_des, stdout_des, ofdes;

    stdout_des = fileno(stdout);
    backup_des = dup(stdout_des);
```



# I/O redirection example (2)

*io-redir.c*

```
printf("Hello, world! (1)\n");

ofdes = open("test.txt", O_WRONLY|O_CREAT|O_TRUNC,
             S_IRUSR|S_IWUSR);

dup2(ofdes, stdout_des);
printf("Hello, world! (2)\n");

dup2(backup_des, stdout_des);
printf("Hello, world! (3)\n");

close(ofdes);
}
```

# Link (1)

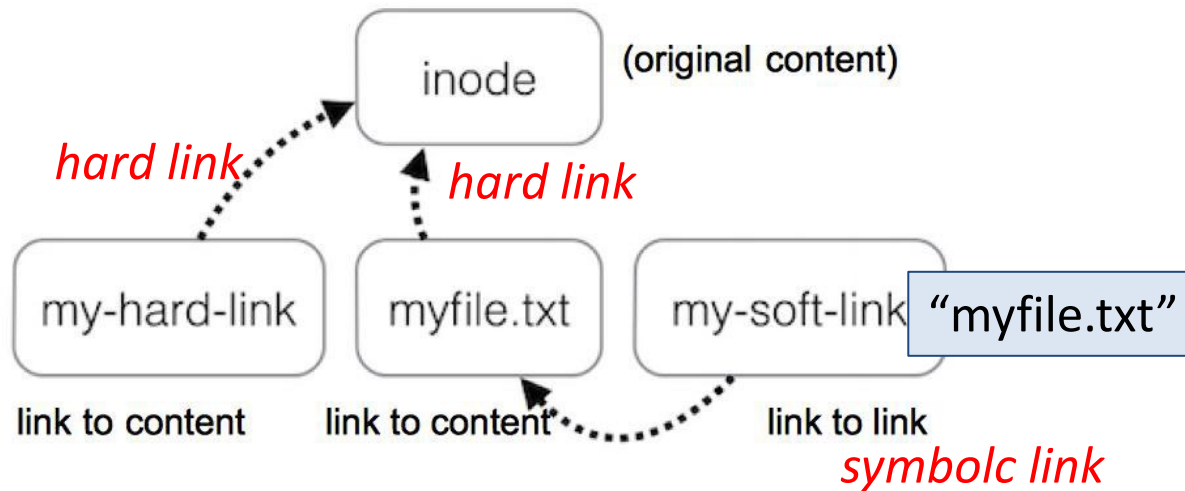
## ■ Symbolic link

- also called *soft link*
- a **file** which records a path name to a target file
- if the target file is removed, the link is not valid any longer, but the symbolic link file still exists

## ■ Hard link

- another link which points to an existing inode which is already used by another file
- an inode can be shared by two or more filenames (hardlinks)
- once a file is update, the update can be seen in all the hardlink files
- though a file is removed, another hardlink can access the file

# Link (2)



## ■ Linux command

- symbolic link:  
`$ ln -s original_file symbolic_link_name`
- hard link:  
`$ ln original_file hard_link_name`

# Hard link

```
#include <unistd.h>
```

```
int link(const char *existing, const char *new_Link);
```

- parameters
  - *existing* : original file name
  - *new\_link* : new link name which will share the inode of existing file
- return
  - 0 if OK
  - -1 on error

# Hard link example (1)

*hlink-ex.c*

```
#include <sys/types.h>
#include <sys/stat.h>
#include <unistd.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    if(argc != 3) {
        perror("argument error");
        return 1;
    }
    if (link(argv[1], argv[2]) < 0) {
        perror("link fail");
        return 2;
    }
}
```

# Hard link example (2)

## ■ *Run & Results*

```
$ ls -l my*
drwxr--r-x 3  oskernel  oskernel  512  Jul 9 21:58  .
-rw-r--r-  1  oskernel  oskernel   15  Jul 9 21:58  myfile

$ ./a.out  myfile  myhardlink

$ ls -l my*
drwxr-xr-x 3  oskernel  oskernel  512  Jul 9 22:01  .
-rw-r--r- 2  oskernel  oskernel   15  Jul 9 22:01  myfile
-rw-r--r- 2  oskernel  oskernel   15  Jul 9 22:01  myhardlink

$
```

# Symbolic link

```
#include <unistd.h>
```

```
int symlink(const char *existing, const char *link_name);
```

- parameters
  - *existing* : original file name
  - *link\_name* : link name which points the existing file
- return
  - 0 if OK
  - -1 on error

# Symbolic link example (1)

*symlink-ex.c*

```
#include <unistd.h>

int main(int argc, char *argv[])
{
    if(argc != 3) {
        perror("argument error");
        return 1;
    }
    if (symlink(argv[1], argv[2]) < 0) {
        perror("symlink fail");
        return 2;
    }
}
```



# Symbolic link example (2)

## ■ *Run & Results*

```
$ ls -l my*  
-rw-r--r-  1 oskernel oskernel  22   Jul 7 22:21  myfile  
  
$ ./a.out myfile mylink  
  
$ ls -l my*  
-rw-r--r--  1 oskernel oskernel  22   Jul 7 22:21  myfile  
lrwxrwxrwx  1 oskernel oskernel   6   Jul 9 22:18  mylink
```

- *What does the file “mylink” contain?*
  - “myfile” → thus, the file size is 6 bytes.

# system calls and symbolic links

system calls which does <b>NOT</b> follow a symbolic link	system calls which follow a symbolic link
lchown, lstat, remove, readlink, rename, unlink  these system calls handle a symbolic link itself as a FILE!	access, chdir, chmod, chown, creat, exec, link, mkdir, mkfifo, mknod, open, opendir, pathconf, stat, truncate

# Following a link (1)

```
#include <unistd.h>
```

```
int readlink(const char *path, void *buf, size_t bufsize);
```

- parameters
  - *path* : link name
  - *buf* : buffer address (original file's name)
  - *bufsize* : buffer size
- return
  - number of bytes read if OK
  - -1 on error

# Following a link example (1)

## *readlink-ex.c*

```
#include <sys/types.h>
#include <sys/stat.h>
#include <unistd.h>
#include <stdio.h>

#define BUFFER_SIZE 100

int main(int argc, char *argv[])
{
    char buf[BUFFER_SIZE];
    int read_size = 0;

    if(argc != 2) {
        perror("argument error");
        return 1;
    }
    if ((read_size = readlink(argv[1], buf, BUFFER_SIZE)) < 0) {
        perror("readlink");
        return 2;
    }

    buf[read_size] = '\0' ;
    printf("%s\n", buf);
}
```

# Following a link example (2)

## *Run & Results*

```
$ ls -ld myfile mylink
-rw-r--r-- 1 oskernel oskernel 22 Jul 7 22:21 myfile
lrwxrwxrwx 1 oskernel oskernel 6 Jul 7 22:18 mylink -> myfile

$ ./a.out mylink
myfile

$
```

# File information retrieval (1)

```
#include<sys/types.h>
#include<sys/stat.h>
#include<unistd.h>

int stat(const char *path, struct stat *buf);
```

- parameters
  - *path* : file path name
  - *buf* : address of struct stat (which contains a file information)
- return
  - 0 if OK
  - -1 on error

# File information retrieval (2)

```
int lstat(const char *path, struct stat *buf);
```

- basically, same as stat(), but
  - if the file is a symbolic link, retrieve the information of the link file itself (does not follow the link)
- parameters
  - *path* : file path name
  - *buf* : address of struct stat (which contains a file information)
- return
  - 0 if OK
  - -1 on error

# File information retrieval (3)

```
int fstat(int fd, struct stat *buf)
```

- parameters
  - *fd* : file descriptor
  - *buf* : address of struct stat (which contains a file information)
- return
  - 0 if OK
  - -1 on error



# struct stat *fields*

```
struct stat {
    dev_t    st_dev;        // device
    ino_t    st_ino;       // i-node #
    mode_t   st_mode;      // access mode
    nlink_t  st_nlink;     // number of hard links
    uid_t    st_uid;       // owner id
    gid_t    st_gid;       // group owner
    dev_t    st_rdev;      // device type (if inode device)
    off_t    st_size;      // total size of file
    long     st_blksize;   // block size for I/O
    long     st_blocks;    // number of blocks allocated
    time_t   st_atime;     // time of last access
    time_t   st_mtime;     // time of last modification
    time_t   st_ctime;     // time of last change (including
                          // ownership change)
};
```

# Macros for struct stat

Macro	Functions
S_ISREG(st_mode)	return true if the file is <i>regular file</i>
S_ISDIR(st_mode)	return true if the file is <i>directory</i>
S_ISCHR(st_mode)	return true if the file is <i>character device file</i>
S_ISBLK(st_mode)	return true if the file is <i>block device file</i>
S_ISFIFO(st_mode)	return true if the file is <i>FIFO file</i>
S_ISLNK(st_mode)	return true if the file is <i>link file</i>
S_ISCOCK(st_mode)	return true if the file is <i>socket file</i>

# File stat example (1)

## *fstat-ex.c*

```
#include <sys/stat.h>
#include <sys/types.h>
#include <unistd.h>
#include <fcntl.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    struct stat statbuf;

    if(argc != 3) {
        perror("argument error");
        return 1;
    }
    if (!strcmp(argv[1], "stat")) {
        if (stat(argv[2], &statbuf) < 0) {
            perror("stat");
            return 2;
        }
    }
}
```

# File stat example (2)

## *fstat-ex.c*

```
    else if (!strcmp(argv[1], "fstat")) {
        int filedes = open(argv[2], O_RDWR);
        if (fstat(filedes, &statbuf) < 0) {
            perror("stat");
            return 3;
        }
    }
    else if (!strcmp(argv[1], "lstat")) {
        if (lstat(argv[2], &statbuf) < 0) {
            perror("lstat");
            return 4;
        }
    }
    if(S_IREG(statbuf.st_mode))
        printf("%s is Regular File\n", argv[2]);
    if(S_ISDIR(statbuf.st_mode))
        printf("%s is Directory\n", argv[2]);
    if(S_ISLNK(statbuf.st_mode))
        printf("%s is Link File\n", argv[2]);
}
```

# File stat example (3)

## ■ Run & Results

```
$ ls -ld mydir myfile mylink
drwxr-xr-x 2    root root  512      Jul 9 19 : 59    mydir
-rw-r--r-- 1    root root   26      Jul 2 23 : 41    myfile
lrwxrwxrwx 1    root root   12      Jul 9 19 : 59    mylink ->mydir/myfile

$ ./a.out stat myfile
myfile is Regular File

$ ./a.out stat mydir
mydir is Directory

$ ./a.out stat mylink
mylink is Regular File

$ ls -l mydir
-rwxrwxrwx 1 peace peace 0 Jul 9 19 : 59 myfile

$ ./a.out lstat mylink
Mylink is Link File

$ ./a.out fstat mylink
mylink is Regular File
```

# Process's Creator (1)

```
#include <sys/types.h>
#include <unistd.h>

uid_t getuid(void)    // process creator's uid
```

- **return**
  - user ID of the process if OK
  - -1 on error

```
#include <sys/types.h>
#include <unistd.h>

uid_t getgid(void)    // process creator's gid
```

- **return**
  - group ID of the process if OK
  - -1 on error

# Process's Creator (2)

```
uid_t geteuid(void)
```

- return

- effective user ID of the process if OK
- -1 on error

- note

- process's effective user id is used as a key for a kernel's protection system, and normally *uid* = *euid*,
- but sometimes *euid* is a different one from *uid* for the dynamic protection system,

```
uid_t getegid(void)
```

- return

- effective group ID of the process if OK
- -1 on error

# Process's Creator (3)

- When a process is created, the user(creator)'s IDs are assigned to the process.
- But, in the following cases, a process's effective IDs are set to a **file owner's** IDs
  - if we run a program file with S\_ISUID (or S\_ISGID) bit, the process's UID is not my UID, but the file owner's UID (or S\_ISGID).

- How to set the bits (example)

```
$ chmod u+s a.out
```

```
$ chmod g+s a.out
```

- Example

```
$ ls -al /bin  
-rwsr-xr-x 1 root root 26492 Dec 1 2017 mount
```

- this **mount** command is executed, the mount process's UID is set to **root**, not the user. → i.e. with the root's authority, the mount will be run!



# Sticky bit: **S\_ISVTX**

- In a directory with sticky bit
  - users can make their own files or subdirectories to the directory
  - but, each file can be only by its owner or supervisor.
- Example

```
$ chmod o+t /test
$ touch /test/file1
$ rm /test/file1 → OK!
$ touch /test/file1
```

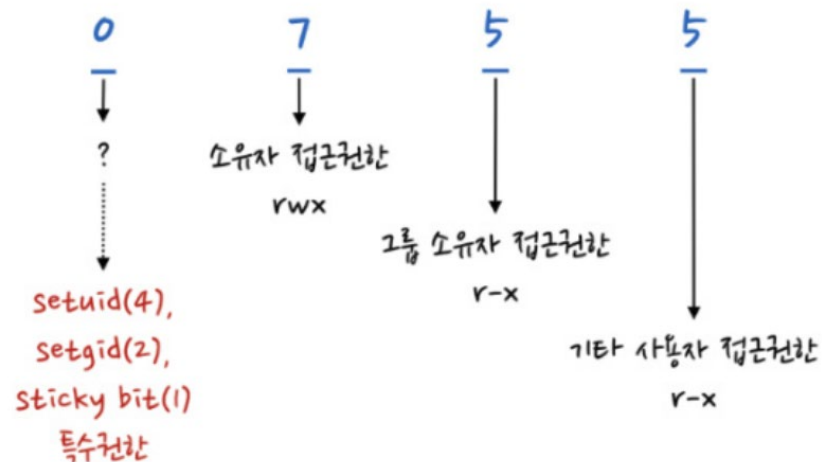
```
....
(another user login)
$ touch /test/file2
$ rm /test/file2 → OK!
$ rm /test/myfile → Failed!
```

# File permission attributes

파일종류	특수권한			소유자 접근권한			그룹 소유자 접근권한			기타 사용자 접근 권한		
-,d,c,b,s,l,p	4	2	1	4	2	1	4	2	1	4	2	1
	setuid	setgid	sticky bit	r	w	x	r	w	x	r	w	x

**d**: directory  
**c**: character device file  
**b**: block device file  
**s**: socket  
**l**: symbolic link

chmod 0755 testfile



# File access of a process

- File access (read/write/execute) is allowed in the following cases
  - if the effective UID of the process is 0 (supervisor)
  - if the effective UID of the process is equal to that of file owner, and if the access permission bit of owner is SET
  - if the effective GID of the process is equal to that of file owner, and if the access permission bit of group is SET
  - if other's access permission bit is SET

# File ownership & IDs

- File owner ID
  - set with the effective UID of the creating process
- File group ID
  - set with the effective GID of the creating process
  - but, if **set-group-ID** bit is set in the creating directory, the file's GID is set to the directory's GID.

# File's access permission

```
#include<unistd.h>
```

```
int access(const char *path, int amode);
```

- check if the process can access a file in the path
- parameter
  - *path* : path name
  - *amode* : access mode for the process to check

amode	meaning
R_OK	<i>READ permission check</i>
W_OK	<i>WRITE permission check</i>
X_OK	<i>Execute or Exploration permission check</i>
F_OK	<i>File existence check</i>

- return
  - 0 if OK
  - -1 on error

# Permission check example

## *access-ex.c*

```
#include<stdio.h>
#include<unistd.h>

int main(int argc, char *argv[])
{
    if(argc < 2) {
        perror("argument error");
        return 1;
    }
    if(access(argv[1], F_OK) == 0) {
        printf("%s : File Exists\n", argv[1]);
    }
    if(access(argv[1], R_OK) == 0)
        printf("%s : Read\n", argv[1]);
    if(access(argv[1], W_OK) == 0)
        printf("%s : Write\n", argv[1]);
    if(access(argv[1], X_OK) == 0)
        printf("%s : Execute\n", argv[1]);
    }
    else printf("%s : NOT exist\n", argv[1]);
}
```

# Default permission change

```
#include<sys/types.h>
#include<sys/stat.h>

mode_t umask(mode_t cmask);
```

- By default
  - by default, a file's permission is set to 0666 (rw-rw-rw)
  - by default, a directory's permission is set to 0777 (rwxrwxrwx)
- umask() changes the default permission
  - set *umask* which masks off (i.e. not permits)
  - e.g. if umask is 0022, a new file permission is set to 0644
- parameter
  - *cmask* : new umask
- return
  - previous umask

```
$ umask
0022
$ touch test
$ ls -al test
-rw-r--r-- root root .....
```

# umask value

- OR'ed combination of these modes

mode	meaning
S_IRWXU	owner RWX
S_IRUSR	owner R
S_IWUSR	owner W
S_IXUSR	owner X
S_IRWXG	group RWX
S_IRGRP	group R
S_IWGRP	group W
S_IXGRP	group X
S_IRWXO	others RWX
S_IROTH	others R
S_IWOTH	others W
S_IXOTH	others X



# File permission change

```
#include<sys/types.h>
#include<sys/stat.h>

int  chmod(const char *path, mode_t mode);
int  fchmod(int fd, mode_t mode);
```

- change the permission mode of a file in the path
  - file owner or supervisor(root) can do this
- parameters
  - *path* : path name of a file
  - *mode* : the new access permission to change
    - mode is also OR'ed combination of the access modes (see p.7)
  - *fd* : file descriptor
- return
  - 0 if OK, -1 on error

# Permission change example (1)

## *chmod-ex.c*

```
#include <sys/stat.h>
#include <sys/types.h>
#include <stdio.h>

int main(int argc, char *argv[])
{
    struct stat statbuf;
    if(argc != 2) {
        perror("argument error");
        return 1;
    }
    if (lstat(argv[1], &statbuf) < 0) {
        perror("lstat");
        return 2;
    }
    if (S_ISREG(statbuf.st_mode)) {
        if(chmod(argv[1], (statbuf.st_mode & ~S_IXGRP)) < 0) {
            perror("chmod");
            return 3;
        }
    }
    else printf("%s is not regular file\n", argv[1]);
}
```

# Permission change example (2)

## ■ Run & Results

```
$ ls -ld myfile mydir
drwx----- 2 root root 512 Jul 15 15 : 13 mydir
-rwx--x--- 1 root root 0    Jul 15 15 : 12 myfile

$ ./a.out myfile
$ ./a.out mydir
mydir is not a regular file

$ ls -ld myfile mydir
drwx----- 2 root root 512 Jul 15 15 : 13 mydir
-rwx----- 1 root root 0    Jul 15 15 : 12 myfile

$
```

# Ownership change

```
#include <unistd.h>
#include <sys/types.h>

int chown(const char *path, uid_t owner, gid_t group);
int lchown(const char *path, uid_t owner, gid_t group);
// doesn't follow links
int fchown(int fd, uid_t owner, gid_t group);
```

- parameters
  - *path* : path name of a file
  - *owner* : owner's UID
  - *group* : group's GID
  - *fd* : file descriptor
- return
  - 0 if OK, -1 on error

# Ownership change example (1)

## *chown-ex.c*

```
#include <sys/types.h>
#include <unistd.h>
#include <stdio.h>
#include <fcntl.h>

int main(int argc, char *argv[])
{
    int owner_id, group_id, filedes ;

    if(argc != 5) {
        perror("argument error");
        return 1;
    }
    owner_id = atoi(argv[3]);
    group_id = atoi(argv[4]);
    if (strcmp(argv[1], "chown") == 0) {
        if (chown(argv[2], owner_id, group_id)) {
            perror("chown");
            return 2;
        }
        printf("chown %s to %s, %s\n", argv[2], argv[3], argv[4]);
    }
}
```

# Ownership change example (2)

## *chown-ex.c*

```
    else if (strcmp(argv[1], "fchown") == 0) {
        filesdes = open(argv[2], O_RDWR);
        if (fchown(argv[2], owner_id, group_id)) {
            perror("chown");
            return 3;
        }
        printf("fchown %s to %s, %s\n", argv[2], argv[3], argv[4]);
    }
    else if (strcmp(argv[1], "lchown") == 0) {
        if (lchown(argv[2], owner_id, group_id)) {
            perror("lchown");
            return 4;
        }
        printf("lchown %s to %s, %s\n", argv[2], argv[3], argv[4]);
    }
}
```

# Ownership change example (3)

## ■ Run & Results

```
$ ls -l my*
drwxr-xr-x 2      oskernel  oskernel  512  Jul 9 21 : 20  mydir/
-rw-r--r- 1      oskernel  oskernel   0   Jul 7 15 : 37  myfile
lrwxrwxrwx 1      oskernel  oskernel  13   Jul 9 21 : 20  mylink -> ~mydir/myfile1

$ id -u cisc
1703

$id -g cisc
511

$ ./a.out chown myfile 1703 511
chown myfile to 1703, 511

$ ls -l my*
drwxr-xr-x 2      oskernel  oskernel  512  Jul 9 21 : 20  mydir/
-rw-r--r-- 1      cisc      cisc       0   Jul 7 15 : 37  myfile
lrwxrwxrwx 1      oskernel  oskernel  13   Jul 9 21 : 20  mylink -> ~mydir/myfile1
```